

RISC-V Assembly Programmer's Manual

Copyright and License Information

The RISC-V Assembly Programmer's Manual is

© 2017 Palmer Dabbelt palmer@dabbelt.com © 2017 Michael Clark michaeljclark@mac.com © 2017 Alex Bradbury asb@lowrisc.org

It is licensed under the Creative Commons Attribution 4.0 International License (CC-BY 4.0). The full license text is available at <https://creativecommons.org/licenses/by/4.0/>.

Command-Line Arguments

I think it's probably better to beef up the binutils documentation rather than duplicating it here.

Registers

Registers are the most important part of any processor. RISC-V defines various types, depending on which extensions are included: The general registers (with the program counter), control registers, floating point registers (F extension), and vector registers (V extension).

General registers

The RV32I base integer ISA includes 32 registers, named `x0` to `x31`. The program counter `PC` is separate from these registers, in contrast to other processors such as the ARM-32. The first register, `x0`, has a special function: Reading it always returns 0 and writes to it are ignored. As we will see later, this allows various tricks and simplifications.

In practice, the programmer doesn't use this notation for the registers. Though `x1` to `x31` are all equally general-use registers as far as the processor is concerned, by convention certain registers are used for special tasks. In assembler, they are given standardized names as part of the RISC-V **application binary interface** (ABI). This is what you will usually see in code listings. If you really want to see the numeric register names, the `-M` argument to `objdump` will provide them.

Register	ABI	Use by convention	Preserved?
x0	zero	hardwired to 0, ignores writes	<i>n/a</i>
x1	ra	return address for jumps	no
x2	sp	stack pointer	yes
x3	gp	global pointer	<i>n/a</i>
x4	tp	thread pointer	<i>n/a</i>

Register	ABI	Use by convention	Preserved?
x5	t0	temporary register 0	no
x6	t1	temporary register 1	no
x7	t2	temporary register 2	no
x8	s0 or fp	saved register 0 or frame pointer	yes
x9	s1	saved register 1	yes
x10	a0	return value or function argument 0	no
x11	a1	return value or function argument 1	no
x12	a2	function argument 2	no
x13	a3	function argument 3	no
x14	a4	function argument 4	no
x15	a5	function argument 5	no
x16	a6	function argument 6	no
x17	a7	function argument 7	no
x18	s2	saved register 2	yes
x19	s3	saved register 3	yes
x20	s4	saved register 4	yes
x21	s5	saved register 5	yes
x22	s6	saved register 6	yes
x23	s7	saved register 6	yes
x24	s8	saved register 8	yes
x25	s9	saved register 9	yes
x26	s10	saved register 10	yes
x27	s11	saved register 11	yes
x28	t3	temporary register 3	no
x29	t4	temporary register 4	no
x30	t5	temporary register 5	no
x31	t6	temporary register 6	no
pc	(none)	program counter	n/a

Registers of the RV32I. Based on RISC-V documentation and Patterson and Waterman "The RISC-V Reader" (2017)

As a general rule, the **saved registers** `s0` to `s11` are preserved across function calls, while the **argument registers** `a0` to `a7` and the **temporary registers** `t0` to `t6` are not. The use of the various specialized registers such as `sp` by convention will be discussed later in more detail.

Control registers

(TBA)

Floating Point registers (RV32F)

(TBA)

Vector registers (RV32V)

(TBA)

Addressing

Addressing formats like `%pcrel_lo()`. We can just link to the RISC-V PS ABI document to describe what the relocations actually do.

Instruction Set

Official Specifications webpage:

- <https://riscv.org/specifications/>

Latest Specifications draft repository:

- <https://github.com/riscv/riscv-isa-manual>

Instructions

RISC-V User Level ISA Specification

<https://riscv.org/specifications/>

RISC-V Privileged ISA Specification

<https://riscv.org/specifications/privileged-isa/>

Instruction Aliases

ALIAS line from `opcodes/riscv-opc.c`

To better diagnose situations where the program flow reaches an unexpected location, you might want to emit there an instruction that's known to trap. You can use an `UNIMP` pseudo-instruction, which should trap in

nearly all systems. The *de facto* standard implementation of this instruction is:

- **C.UNIMP: 0000**. The all-zeroes pattern is not a valid instruction. Any system which traps on invalid instructions will thus trap on this **UNIMP** instruction form. Despite not being a valid instruction, it still fits the 16-bit (compressed) instruction format, and so **0000 0000** is interpreted as being two 16-bit **UNIMP** instructions.
- **UNIMP : C0001073**. This is an alias for **CSRRW x0, cycle, x0**. Since **cycle** is a read-only CSR, then (whether this CSR exists or not) an attempt to write into it will generate an illegal instruction exception. This 32-bit form of **UNIMP** is emitted when targeting a system without the C extension, or when the **.option norvc** directive is used.

Pseudo Ops

Both the RISC-V-specific and GNU **.-**prefixed options.

The following table lists assembler directives:

Directive	Arguments	Description
.align	integer	align to power of 2 (alias for .p2align)
.file	"filename"	emit filename FILE LOCAL symbol table
.globl	symbol_name	emit symbol_name to symbol table (scope GLOBAL)
.local	symbol_name	emit symbol_name to symbol table (scope LOCAL)
.comm	symbol_name,size,align	emit common object to .bss section
.common	symbol_name,size,align	emit common object to .bss section
.ident	"string"	accepted for source compatibility
.section	[{.text,.data,.rodata,.bss}]	emit section (if not present, default .text) and make current
.size	symbol, symbol	accepted for source compatibility
.text		emit .text section (if not present) and make current
.data		emit .data section (if not present) and make current
.rodata		emit .rodata section (if not present) and make current
.bss		emit .bss section (if not present) and make current
.string	"string"	emit string
.asciz	"string"	emit string (alias for .string)
.equ	name, value	constant definition
.macro	name arg1 [, argn]	begin macro definition \argname to substitute
.endm		end macro definition

Directive	Arguments	Description
.type	symbol, @function	accepted for source compatibility
.option	{rvc,norvc,pic,nopic,push,pop}	RISC-V options
.byte	expression [, expression]*	8-bit comma separated words
.2byte	expression [, expression]*	16-bit comma separated words
.half	expression [, expression]*	16-bit comma separated words
.short	expression [, expression]*	16-bit comma separated words
.4byte	expression [, expression]*	32-bit comma separated words
.word	expression [, expression]*	32-bit comma separated words
.long	expression [, expression]*	32-bit comma separated words
.8byte	expression [, expression]*	64-bit comma separated words
.dword	expression [, expression]*	64-bit comma separated words
.quad	expression [, expression]*	64-bit comma separated words
.dtprelword	expression [, expression]*	32-bit thread local word
.dtpreldword	expression [, expression]*	64-bit thread local word
.sleb128	expression	signed little endian base 128, DWARF
.uleb128	expression	unsigned little endian base 128, DWARF
.p2align	p2,[pad_val=0],max	align to power of 2
.balign	b,[pad_val=0]	byte align
.zero	integer	zero bytes

Assembler Relocation Functions

The following table lists assembler relocation expansions:

Assembler Notation	Description	Instruction / Macro
%hi(symbol)	Absolute (HI20)	lui
%lo(symbol)	Absolute (LO12)	load, store, add
%pcrel_hi(symbol)	PC-relative (HI20)	auipc
%pcrel_lo(label)	PC-relative (LO12)	load, store, add
%tprel_hi(symbol)	TLS LE "Local Exec"	lui
%tprel_lo(symbol)	TLS LE "Local Exec"	load, store, add
%tprel_add(symbol)	TLS LE "Local Exec"	add

Assembler Notation	Description	Instruction / Macro
<code>%tls_ie_pcrel_hi(symbol) *</code>	TLS IE "Initial Exec" (HI20)	<code>auipc</code>
<code>%tls_gd_pcrel_hi(symbol) *</code>	TLS GD "Global Dynamic" (HI20)	<code>auipc</code>
<code>%got_pcrel_hi(symbol) *</code>	GOT PC-relative (HI20)	<code>auipc</code>

* These reuse `%pcrel_lo(label)` for their lower half

Labels

Text labels are used as branch, unconditional jump targets and symbol offsets. Text labels are added to the symbol table of the compiled module.

```
loop:
    j loop
```

Numeric labels are used for local references. References to local labels are suffixed with 'f' for a forward reference or 'b' for a backwards reference.

```
1:
    j 1b
```

Absolute addressing

The following example shows how to load an absolute address:

```
.section .text
.globl _start
_start:
    lui a0,      %hi(msg)      # load msg(hi)
    addi a0, a0,  %lo(msg)      # load msg(lo)
    jal ra, puts
2:      j 2b

.section .rodata
msg:
    .string "Hello World\n"
```

which generates the following assembler output and relocations as seen by `objdump`:

```
0000000000000000 <_start>:
  0:  000005b7                lui      a1,0x0
                        0:  R_RISCV_HI20 msg
```

```

4:      00858593                addi      a1,a1,8 # 8 <.L21>
                                4: R_RISCV_L012_I      msg

```

Relative addressing

The following example shows how to load a PC-relative address:

```

.section .text
.globl _start
_start:
1:      auipc a0,      %pcrel_hi(msg) # load msg(hi)
        addi a0, a0, %pcrel_lo(1b) # load msg(lo)
        jal ra, puts
2:      j 2b

.section .rodata
msg:
        .string "Hello World\n"

```

which generates the following assembler output and relocations as seen by objdump:

```

0000000000000000 <_start>:
 0:      00000597                auipc      a1,0x0
                                0: R_RISCV_PCREL_HI20      msg
 4:      00858593                addi      a1,a1,8 # 8 <.L21>
                                4: R_RISCV_PCREL_L012_I   .L11

```

GOT-indirect addressing

The following example shows how to load an address from the GOT:

```

.section .text
.globl _start
_start:
1:      auipc a0, %got_pcrel_hi(msg) # load msg(hi)
        ld    a0, %pcrel_lo(1b)(a0) # load msg(lo)
        jal ra, puts
2:      j 2b

.section .rodata
msg:
        .string "Hello World\n"

```

which generates the following assembler output and relocations as seen by objdump:

```

0000000000000000 <_start>:
    0:    00000517                auipc    a0,0x0
                                0: R_RISCV_GOT_HI20      msg
    4:    00053503                ld      a0,0(a0) # 0 <_start>
                                4: R_RISCV_PCREL_L012_I  .L11

```

Load Immediate

The following example shows the **li** psuedo instruction which is used to load immediate values:

```

.section .text
.globl _start
_start:

.equ CONSTANT, 0xcafebabe

    li a0, CONSTANT

```

which generates the following assembler output as seen by objdump:

```

0000000000000000 <_start>:
    0:    00032537                lui      a0,0x32
    4:    bfb50513                addi     a0,a0,-1029
    8:    00e51513                slli     a0,a0,0xe
   c:    abe50513                addi     a0,a0,-1346

```

Load Address

The following example shows the **la** psuedo instruction which is used to load symbol addresses:

```

.section .text
.globl _start
_start:

    la a0, msg

.section .rodata
msg:
    .string "Hello World\n"

```

which generates the following assembler output and relocations for non-PIC as seen by objdump:

```

0000000000000000 <_start>:
    0:    00000517                auipc    a0,0x0

```



```

                                0: R_RISCV_PCREL_HI20    msg
4:  00850513                    addi    a0,a0,8 # 8 <_start+0x8>
                                4: R_RISCV_PCREL_LO12_I    .L11

```

and generates the following assembler output and relocations for PIC as seen by objdump:

```

0000000000000000 <_start>:
    0:  00000517                    auipc    a0,0x0
                                0: R_RISCV_GOT_HI20    msg
    4:  00053503                    ld      a0,0(a0) # 0 <_start>
                                4: R_RISCV_PCREL_LO12_I    .L0

```

Constants

The following example shows loading a constant using the %hi and %lo assembler functions.

```

.equ UART_BASE, 0x40003000

    lui a0,    %hi(UART_BASE)
    addi a0, a0, %lo(UART_BASE)

```

This example uses the `li` pseudoinstruction to load a constant and writes a string using polled IO to a UART:

```

.equ UART_BASE, 0x40003000
.equ REG_RBR, 0
.equ REG_TBR, 0
.equ REG_IIR, 2
.equ IIR_TX_RDY, 2
.equ IIR_RX_RDY, 4

.section .text
.globl _start
_start:
1:      auipc a0, %pcrel_hi(msg)    # load msg(hi)
        addi a0, a0, %pcrel_lo(1b) # load msg(lo)
2:      jal ra, puts
3:      j 3b

puts:
        li a2, UART_BASE
1:      lbu a1, (a0)
        beqz a1, 3f
2:      lbu a3, REG_IIR(a2)
        andi a3, a3, IIR_TX_RDY
        beqz a3, 2b
        sb a1, REG_TBR(a2)
        addi a0, a0, 1

```

```

        j 1b
3:      ret

.section .rodata
msg:
        .string "Hello World\n"

```

Floating-point rounding modes

For floating-point instructions with a rounding mode field, the rounding mode can be specified by adding an additional operand. e.g. `fcvt.w.s` with round-to-zero can be written as `fcvt.w.s a0, fa0, rtz`. If unspecified, the default `dyn` rounding mode will be used.

Supported rounding modes are as follows (must be specified in lowercase):

- `rne`: round to nearest, ties to even
- `rtz`: round towards zero
- `rdn`: round down
- `rup`: round up
- `rmn`: round to nearest, ties to max magnitude
- `dyn`: dynamic rounding mode (the rounding mode specified in the `frm` field of the `fcsr` register is used)

Control and Status Registers

The following code sample shows how to enable timer interrupts, set and wait for a timer interrupt to occur:

```

.equ RTC_BASE,      0x40000000
.equ TIMER_BASE,    0x40004000

# setup machine trap vector
1:    auipc    t0, %pcrel_hi(mtvec)      # load mtvec(hi)
      addi     t0, t0, %pcrel_lo(1b)     # load mtvec(lo)
      csrrw    zero, mtvec, t0

# set mstatus.MIE=1 (enable M mode interrupt)
      li       t0, 8
      csrrs    zero, mstatus, t0

# set mie.MTIE=1 (enable M mode timer interrupts)
      li       t0, 128
      csrrs    zero, mie, t0

# read from mtime
      li       a0, RTC_BASE
      ld       a1, 0(a0)

# write to mtimecmp
      li       a0, TIMER_BASE
      li       t0, 1000000000

```

```
        add    a1, a1, t0
        sd     a1, 0(a0)

# loop
loop:
        wfi
        j loop

# break on interrupt
mtvec:
        csrrc  t0, mcause, zero
        bgez  t0, fail      # interrupt causes are less than zero
        slli  t0, t0, 1     # shift off high bit
        srli  t0, t0, 1
        li    t1, 7        # check this is an m_timer interrupt
        bne   t0, t1, fail
        j     pass

pass:
        la     a0, pass_msg
        jal    puts
        j     shutdown

fail:
        la     a0, fail_msg
        jal    puts
        j     shutdown

.section .rodata

pass_msg:
        .string "PASS\n"

fail_msg:
        .string "FAIL\n"
```